

Kevin Chen

865-235-0680 | kevin.chen.kc2593@yale.edu | <https://www.linkedin.com/in/kevinchenpage>

SUMMARY STATEMENT

Policy professional and technical researcher with 3+ years of experience in high-security environments at the CIA and FBI, and over 7+ years of programming experience. Focused on the intersection of AI assurance, technical governance, and international security. Currently pursuing an MPP at Yale with a research focus on technical CBRN guardrailings, AI-assisted misinformation campaigns, and great power competition between the US and China.

EDUCATION

Yale University

New Haven, CT

Master in Public Policy in Global Affairs

Aug. 2025 – June 2027

- Awarded the Nelson Cunningham '80 Fellowship Fund (Full-tuition merit scholarship).
- Current Editor-in-Chief for Yale Journal of International Affairs.
- Graduate Teaching Fellow for Prof. Edward Wittenstein's course, *Artificial Intelligence, Emerging Technologies, and National Power*.

Sciences Po Paris

Paris, France

Stablecoins, DLT, and International Trade (Study Abroad)

March 2026 – April 2026

- Collaborated with the International Chamber of Commerce (ICC) to develop stablecoin and distributed ledger technology (DLT) frameworks.

Dartmouth College

Hanover, NH

Bachelor of Arts, Quantitative Social Science, Honors

Sept. 2018 – June 2022

EXPERIENCE

Centre for the Governance of AI (GovAI)

London, UK

Summer Research Fellow

June 2026 – Aug. 2026

- Constructing an empirical evaluation framework simulating fixed-horizon negotiation protocols (10-turn loops) on local workstations to quantify auditor model detectability curves as a strict function of turn horizon.
- Engineering an automated benchmarking pipeline across Western frontier models and regional open-weight architectures (Qwen, DeepSeek, Mistral) utilizing a curated question bank of 100+ specific geopolitical flashpoints.
- Evaluating the structural failure modes of post-hoc log auditing against active auditing paradigms, mapping behavioral anomalies and tracking internal model reasoning via hidden chain-of-thought (CoT) protocols.
- Targeting technical research outputs for the NeurIPS AI Safety workshop alongside specialized policy memos designed for agentic risk evaluation teams at the US and UK AI Safety Institutes (AISIs).

Central Intelligence Agency – Directorate of Digital Innovation

Langley, VA

Data Scientist – AI/ML Engineer – TS/SCI Full Scope Polygraph

Oct. 2023 – May 2025

- Implemented AI assurance frameworks and adversarial testing for Mandarin-language models to evaluate model alignment and robustness.
- Built and deployed NLP and computer vision pipelines using Python and PyTorch for mission-critical applications.
- Led cross-functional teams to deliver enterprise AI features under strict security and performance requirements.
- Designed scalable ETL pipelines across multilingual datasets using Kubernetes and Docker.

Federal Bureau of Investigation – Resource Planning Office

Washington, D.C.

Data Analyst – Resource Analysis – TS/SCI Full Scope Polygraph

Aug. 2022 – Sept. 2023

- Engineered SQL-based pipelines to process personnel data for 30,000+ employees, reducing reporting latency.
- Built interactive data visualizations in R and Python to inform multi-million-dollar budget decisions.
- Innovated projects to process massive data volumes for FBI positions and quarterly personnel expense requests.

LEADERSHIP AND ACADEMIC RESEARCH

Supervised Program for Alignment Research (SPAR) – Kairos

Remote

Project Mentor

Sept. 2026 – Dec. 2026

- Selected to co-mentor *Assessing Taiwan's Global Position for Transformative AI* for the Fall 2026 SPAR cohort, a competitively reviewed alignment research program administered by Kairos.
- Structuring the project into discrete research modules with individual mentee ownership, and conducting rolling application review, technical testing, and interviews to select the research team.
- Providing weekly research supervision and facilitating midterm and final report milestones assessing Taiwan's institutional and industrial preparedness for transformative AI.

California Governor's Office of Emergency Services (Cal OES) Sacramento, CA (Remote)

Research Collaborator

June 2026 – Present

- Working with State AI Science Advisor Michael Chen to operationalize California's SB 53 AI governance framework into practical public-sector crisis management and emergency response playbooks.
- Drafting technical protocols and threshold metrics defining catastrophic "loss of control" incidents, translating complex frontier model alignment failure modes into actionable state-level intervention triggers.
- Developing policy recommendations for state-level compute and deployment guardrails, focusing on technical verification mechanisms and mandatory model evaluation disclosure requirements for California-based frontier labs.

U.S. Department of State Honolulu, HI

Selected Participant, Young Trilateral Leaders (YTL) Summit

Aug. 2026

- Selected by the U.S. Mission to participate in the flagship trilateral summit advancing U.S.-Japan-ROK technology policy and multilateral diplomacy in the Indo-Pacific.
- Collaborate in cross-national teams with American, Korean, and Japanese policymakers to engineer actionable project proposals leveraging AI to reinforce regional stability and economic security.

OpenAI San Francisco, CA (Remote)

Research Collaborator

May 2026 – Present

- Co-authoring a policy research paper with Ziad Reslan and Kim Malfacini of OpenAI evaluating global frameworks for advanced AI access and hardware/compute governance.
- Develop structural outlines and policy recommendations addressing the trade-offs between open-source model dissemination, technical verification mechanisms, and international safety guardrails.

Council on Foreign Relations New York City, NY (Remote)

Graduate Research Assistant to Adjunct Senior Fellow Gordon Goldstein

May 2026 – Present

- Conduct targeted literature reviews and structural analyses for a forthcoming Rockefeller Brothers Fund/Council on Foreign Relations book project on artificial intelligence and international security by Gordon Goldstein.
- Investigate geopolitically critical facets of US-China technology competition, evaluating the strategic implications of advanced computing capabilities on regional stability and international safety frameworks.

Doublethink Lab Taipei, Taiwan (Hybrid)

Research Collaborator

May 2026 – Present

- Developing a year-long Master's thesis titled in direct partnership with Doublethink Lab researchers to investigate AI-assisted cognitive warfare across the Taiwan Strait.
- Engineer quantitative frameworks to detect and analyze Foreign Information Manipulation and Interference (FIMI) campaigns and utilize multi-lingual natural language processing to identify coordinated inauthentic behavior.
- Coordinating an in-person data-sprint and research workshop in Taipei for October 2026 to cross-validate machine learning threat models with local behavioral datasets.

PUBLICATIONS

"Domestic Compute as a Critical Chokepoint" | *Department of Homeland Security* June 2026

- Technical white paper for DHS NCITE analyzing frontier AI data centers and domestic compute infrastructure as asymmetric operational targets.

"Geospatial AI Frameworks for Infrastructure Monitoring" | *TechNatSec Journal* Feb. 2026

- National winner for paper competition with the MIT-Harvard Technology & National Security Conference.
- Applied multi-modal satellite imagery and Computer Vision (CNNs) to identify structural damage in conflict zones.
- Fused GeoAI outputs with ACLED and WFP telemetry to generate predictive risk indicators for civilian harm.

"Vertical integration of rare Earth elements for US autonomous dominion" | *CNAS* Nov. 2025

- Winning essay on building a domestic rare earth elements supply chain for the Center for a New American Security (CNAS) Pathfinder research competition.

TECHNICAL SKILLS

Programming Languages: C++, Java, Python, R, SQL, LaTeX

Libraries: Tensorflow, Pytorch, Scikit-learn, XGBoost, Pandas, NumPy, SpaCy, tidyxl

Databases/Engineering: Elasticsearch, Kafka, MySQL, PostgreSQL, Kubernetes, Triton, Docker

AI Tools: Hugging Face Transformers, YOLOv8, Detectron2, MMDetection, LangChain, NLTK

Spoken Languages: Mandarin Chinese (Fluent), Russian (Proficient)

EXTRACURRICULARS AND AWARDS

Technical AI Safety Certificate — <i>Bluedot Impact</i>	May. 2026
TechnatSec Symposium Finalist — <i>Harvard Belfer Center and MIT Sloan</i>	Mar. 2026
CNAS Pathfinder Essay National Winner — <i>Center for a New American Security</i>	Nov. 2025
Asian Pacific American Organization Directorate Representative — <i>CIA</i>	Oct. 2024
Pat Roberts Intelligence Scholars Program — <i>CIA</i>	June 2024
Intermediate Data Exploitation Certification — <i>FBI</i>	Jan. 2023